

Contents

FinTechCo Deck — Full Audit & Recommendations	1
Overall verdict	1
Scorecard	1
The title problem, in one sentence	2
Slide-by-slide audit	3
Cross-cutting recommendations	7
Proposed title set (copy-paste ready)	8
Priority order, if you only have an hour	10

FinTechCo Deck — Full Audit & Recommendations

Audit target: FinTechCo_Claude_Code_Demo_v3.pptx (20 slides) Audience: Final-round Anthropic Solutions Architect / Solutions Engineer interview panel Author: Holden Otolini

Overall verdict

Strong bones. The tone has swung too far toward “soft.”

The v3 deck successfully fixed the aggressive, bro-y tone of v2. But in removing dashes and softening language, we introduced a second problem: the deck now reads slightly apologetic in places (“easy,” “gentle,” “small asks,” “friendly walk through”). That hurts you against Anthropic’s SA rubric.

Based on 2026 candidate reports, the Anthropic SA/SE bar specifically looks for:

1. **Measured confidence.** You’re the expert in the room on Claude Code. Sound like it.
2. **Customer problem framing before product.** You’re already doing this well.
3. **Honest tradeoffs, not polished pitches.** The limitations slide is your single biggest asset — move it earlier.
4. **Systems thinking.** Talk about queues, scale, failure modes, and “what breaks at 180 engineers.”
5. **Simplicity over cleverness.** Anthropic’s own value: “*we don’t invent a spaceship if all we need is a bicycle.*” Your deck should feel like a bicycle.
6. **Authenticity. Limited small talk. High bar for directness.**

The fix is not to un-do the friendliness. It’s to re-introduce *professional conviction* in the titles while keeping the warm body copy. Think “trusted advisor” not “eager vendor.”

Scorecard

Dimension	Score	Notes
Narrative arc	8 / 10	Agenda → Problem → Discovery → Demo → ROI → Rollout → Close. Textbook.
Title clarity	5 / 10	Several titles are subtitles pretending to be titles. See rewrites below.
Tone calibration	6 / 10	Leaned too far toward “gentle.” Should be <i>warm</i> + <i>confident</i> , not <i>apologetic</i> .
Executive scannability	7 / 10	If the CTO reads only the titles, do they get the story? Today: not quite.
Business framing	8 / 10	ROI and rollout are strong. Pricing math lands.
Technical credibility	7 / 10	Architecture is implicit. One slide should make the flow visible (Safety Gate, CLAUDE.md, hooks).
Honesty / trade-offs	9 / 10	Slide 19 is your secret weapon. Anthropic specifically looks for this.
Anthropic values alignment	7 / 10	Safety is mentioned but not framed as <i>Anthropic’s identity</i> .

The title problem, in one sentence

Most of your slide titles are currently the *subtitle*, and the real title is buried in small-caps above it. On a projector, the first thing the room reads is the headline. Right now the headline is doing all the narrative lifting with zero help from the eyebrow label.

Rule for every slide: the room should be able to read only the main title and still know what this slide is *for*. Treat the small-caps eyebrow as *section*; treat the title as *takeaway*.

Slide-by-slide audit

Slide 1 — Cover

Current: Eyebrow: *“A learning session with FinTechCo”*. Title: *“Claude Code”*. Subtitle: *“A friendly walk through how your engineers, SREs, and data scientists can ship faster, together.”*

Issues: - *“Learning session”* sounds like training, not a pitch. You are selling, not teaching. - *“Friendly walk through”* undersells you.

Recommendation: - Eyebrow → A working session with FinTechCo (neutral, professional) - Title → Claude Code (keep) - Subtitle → How 180 engineers ship faster, with the controls your auditors expect.

Why: gives the CTO the two things they care about in one line — velocity and control.

Slide 2 — Agenda

Current title: *“Here’s how I’d love to spend our next 40 minutes together.”*

Issues: - Reads as a sentence, not a title. 11 words before the room learns this is the agenda. - *“I’d love to”* is weaker than *“I’ll.”* Solutions Architects *drive*; they don’t *ask*.

Recommendation: - Eyebrow → AGENDA · 40 MINUTES - Title → Here's how we'll use our time together. - Keep the item descriptions exactly as-is — those are strong. - Item 03 *“Discovery session”* body copy: change *“Five easy questions”* → *“Five questions”*. *“Easy”* is unnecessary and slightly diminishing.

Slide 3 — What is Claude Code

Current title: *“Think of it as a teammate in the terminal.”*

Issues: - Cute but doesn’t answer the question the eyebrow asks. The room wants a definition before a metaphor. - The three columns (*“It reads your whole codebase / It plans before it acts / It checks its own work”*) are good. Keep.

Recommendation: - Eyebrow → WHAT IS CLAUDE CODE - Title → A coding agent that works the way a senior engineer does. - Keep the supporting paragraph.

Why: Anthropic SA candidates are evaluated on whether they can explain AI capability to non-technical stakeholders. *“Coding agent”* is the accurate term; *“works the way a senior engineer does”* is the accessible frame.

Slide 4 — Why are we here today

Current title: *“Before I share anything, here’s what I think I’m hearing.”*

Issues: - Great instinct (demonstrating active listening). Bad title (too long, buries the takeaway). - The four cards are excellent.

Recommendation: - Eyebrow → WHAT I'M HEARING - Title → Four things on FinTechCo's plate, from what you've shared so far. - Subtitle (keep, tighten) → Correct me where I'm wrong. This is the most valuable thirty seconds of the call.

Slide 5 — Discovery session

Current title: *"Five easy questions. Your answers will shape the demo."*

Issues: - "Easy" is doing you no favors. Take it out. - Q1-Q5 are well-scoped. Keep.

Recommendation: - Eyebrow → DISCOVERY - Title → Five questions before I show you anything. - Keep the closing line about playback — that's a senior move.

Slide 6 — Three teams at FinTechCo

Current title: *"Three teams, three very different problems."*

Issues: Title is fine. The "120 / 20 / 40" headcounts with TODAY / WITH CLAUDE CODE comparison is the single strongest slide in the deck. Don't touch it.

Recommendation: - Eyebrow → WHO THIS HELPS AT FINTECHCO (keep) - Title → Three teams. Three different bottlenecks. One tool. (slightly tighter)

Slide 7 — Safety, security, governance

Current title: *"Built to feel comfortable in the environment you already work in."*

Issues: - "Feel comfortable" is hedged language. For a regulated fintech audience, this is where you stand firm. - Compliance strip is well-placed.

Recommendation: - Eyebrow → SAFETY, SECURITY, GOVERNANCE - Title → Designed for a regulated environment from the start. - Keep all four supporting pillars.

Slide 8 — Live demo opener

Current title: *"Let's walk through a morning with Claude Code."*

Issues: Good. Keep.

Recommendation: no change, except consider bolding slide 8 as the *center of gravity* of the deck. Every slide before it builds toward it; every slide after it refers back to it.

Slide 9 — What you just saw

Current title: *"Twelve minutes, four things you'd notice on Monday morning."*

Issues: Strong. Keep.

Recommendation: no change.

Slide 10 — ROI framework

Current title: *"A framework, not a promise. We'll use your numbers."*

Issues: - Excellent title. Do not change. - "about 50% / 60% / 60% / 70%" is correct calibration — specific enough to be credible, hedged enough to be honest. - The math callout (\$1,500 savings vs \$100 cost) is the single best ROI slide in the deck. Leave it alone.

Slide 11 — Rollout plan

Current title: *"Three weeks to a signal. Three months to a decision."*

Issues: One of the strongest titles in the deck. Don't touch it.

Recommendation: no change.

Slide 12 — Next steps

Current title: *"Three small asks. None of them require procurement today."*

Issues: - "Small asks" is slightly diminutive for a CTO-level conversation. You *are* asking for real things (access to legal, a named owner, a 30-minute setup call, and a pilot readout). - The "No procurement today" framing is excellent and should stay.

Recommendation: - Eyebrow → NEXT STEPS - Title → Three asks. None of them need procurement today.

Slide 13 — Appendix divider

Current title: *"A few things you may want to dig into later in our conversation."*

Issues: Too long, too hedged.

Recommendation: - Title → Appendix: the questions your team will ask next. - Subtitle (list): keep as-is.

Slide 14 — Why Anthropic

Current title: *"This is as much a company question as it is a product question."*

Issues: - Abstract. The room doesn't know what claim you're making. - Content is strong (Constitutional AI, no training, 200k context, SOC 2, Claude plans before it acts). Title should preview the point.

Recommendation: - Eyebrow → WHY ANTHROPIC, WHY CLAUDE - Title → You're not just buying a tool. You're picking a partner your auditors will meet.

Why: this is the argument that wins against Copilot, Cursor, and Devin in a regulated-industry conversation. Say it out loud.

Slide 15 — How we compare

Current title: *"Different tools for different problems."*

Issues: OK, but generic. You're making a sharper point inside the slide than the title suggests.

Recommendation: - Eyebrow → COMPETITIVE LANDSCAPE - Title → The question isn't which AI tool. It's which problem you're solving. - (That sentence is already in your body copy. Promote it to the headline.)

Slide 16 — Data, security, and your auditors

Current title: *"Plain answers to the questions you'll hear first."*

Issues: Strong. Keep.

Recommendation: no change.

Slide 17 — CLAUDE.md at scale

Current title: *"How 180 engineers get the same rules on day one."*

Issues: Strong. Keep.

Recommendation: no change.

Slide 18 — MCP integration

Current title: *"Claude Code connects to the tools your team already uses."*

Issues: Fine. Could be sharper.

Recommendation: - Title → Claude Code plugs into Jira, Confluence, and your internal APIs.

Why: specificity builds credibility. A CTO scanning titles wants to see Jira and Confluence called out by name.

Slide 19 — Honest limitations

Current title: *“Being honest about this makes everything else more credible.”*

Issues: - This is the most important slide in the whole deck for an Anthropic interview. The candidate research is unanimous: they evaluate *“candidates who confidently present a textbook solution and resist exploring alternatives miss the point. The interviewer wants to see you reason through trade-offs.”* - Your title is meta (talking about why you’re being honest) instead of framing the honesty itself.

Recommendation: - Eyebrow → HONEST LIMITATIONS - Title → Five things Claude Code won't do, so you can plan around them.

Also: consider moving this slide *earlier* — before the ROI slide — if the panel is senior enough. It disarms skepticism and earns you the right to make the ROI claim that follows.

Slide 20 — Closing

Current: *“Let’s build the pilot together.”* with *“I’d love to end our time with a name and a date, rather than a thank you.”*

Issues: Almost perfect. One small tune-up.

Recommendation: - Title → Let's build the pilot together. (keep) - Subtitle → I'd rather end on a name and a date than a thank-you. So: who on your side is the natural owner?

Why: compressing “I’d love to end our time with...” → “I’d rather end on...” lands harder. Ending on the question forces the room to answer you.

Cross-cutting recommendations


1. Re-balance the tone (most important)

Audit every instance of these hedging words and decide case-by-case whether to keep them: *easy, small, gentle, friendly, I’d love to, hopefully, just, a little, kind of, sort of, feel comfortable, tends to*.

Rule of thumb: keep hedges in *body copy*, remove them from *titles*. Titles should sound like a senior person saying “here’s what I know.” Body copy can be the warmer human voice.

2. Add one architecture visual

The deck has zero systems diagrams. For a fintech CTO, one clean box-and-arrow diagram showing:



would carry more technical credibility than any paragraph. Drop it in at slide 7 or 17.

3. Add a “what would break” slide (optional but high-leverage)

Anthropic’s SA rubric explicitly asks: *where does this fail at scale?* You already have slide 19 (what Claude doesn’t do). Consider a companion slide covering *failure modes*: “what happens when Claude is wrong / when CLAUDE.md is wrong / when a pilot engineer disagrees with a proposed change.” This is the single most common missing piece in Solutions Architect demos.

4. Strengthen the Anthropic-mission thread

Per 2026 candidate reports: “Anthropic weaves its mission of AI safety into every fiber of its hiring process.” Your deck mentions SOC 2 and Constitutional AI but doesn’t explicitly connect Claude Code’s behavior (plans before acting, Safety Gate, honest-about-limits) to Anthropic’s identity as a safety-first lab. A single line on slide 7 or 14 —“the Safety Gate isn’t a feature we bolted on; it’s how Anthropic believes agentic work should behave”— makes that thread visible to the panel.

5. Run the title-only scan

Before the interview, print the deck six-up, titles only. If a non-technical stakeholder can follow the narrative from titles alone, you’re ready. If they can’t, the titles are still too decorative.

Proposed title set (copy-paste ready)

#	Eyebrow	Title
1	A working session with FinTechCo	Claude Code
2	AGENDA · 40 MINUTES	Here’s how we’ll use our time together.
3	WHAT IS CLAUDE CODE	A coding agent that works the way a senior engineer does.
4	WHAT I’M HEARING	Four things on FinTechCo’s plate, from what you’ve shared so far.

#	Eyebrow	Title
5	DISCOVERY	Five questions before I show you anything.
6	WHO THIS HELPS AT FINTECHCO	Three teams. Three different bottlenecks. One tool.
7	SAFETY, SECURITY, GOVERNANCE	Designed for a regulated environment from the start.
8	LIVE DEMO	Let's walk through a morning with Claude Code.
9	WHAT YOU JUST SAW	Twelve minutes, four things you'd notice on Monday morning.
10	MEASURING THE RETURN	A framework, not a promise. We'll use your numbers.
11	A PRAGMATIC ROLLOUT	Three weeks to a signal. Three months to a decision.
12	NEXT STEPS	Three asks. None of them need procurement today.
13	APPENDIX	The questions your team will ask next.
14	WHY ANTHROPIC, WHY CLAUDE	You're picking a partner your auditors will meet.
15	COMPETITIVE LANDSCAPE	The question isn't which AI tool. It's which problem you're solving.
16	DATA, SECURITY, AUDITORS	Plain answers to the questions you'll hear first.
17	CLAUDE.MD · GOVERNANCE AT SCALE	How 180 engineers get the same rules on day one.
18	MCP · INTEGRATION	Claude Code plugs into Jira, Confluence, and your internal APIs.
19	HONEST LIMITATIONS	Five things Claude Code won't do, so you can plan around them.
20	THANK YOU	Let's build the pilot together.

Priority order, if you only have an hour

1. **Update the 9 titles flagged in bold above** (slides 1, 2, 3, 4, 5, 7, 12, 14, 15, 18, 19). Biggest signal-to-effort ratio.
2. **Sweep body copy for “easy,” “small,” “gentle,” “I’d love to”** — keep where warm, cut where weak.
3. **Add the architecture diagram** to slide 7 or 17.
4. **Rehearse slide 19 (limitations)** as a confident moment, not an apologetic one. Anthropic panels *reward* this specifically.

Everything else in the deck is already working.