

Contents

Claude Code for FinTechCo	1
The problem we’re solving	1
What Claude Code does differently	1
The numbers (conservative)	1
Security posture	2
A pragmatic rollout	2
What we need from you to start	3
Honest limitations	3
Who’s your partner	3

Claude Code for FinTechCo

A coding agent built for regulated environments. Safety at the model level, governance in your repo, a human on every write.

The problem we’re solving

FinTechCo runs 180 engineers under PCI-DSS and FFIEC oversight. Three things compound: new engineers take a week to safely touch unfamiliar code; P1 incidents lose 45 minutes between alert and root cause; and pressure to ship faster keeps rising without the controls catching up.

What Claude Code does differently

Reads the whole codebase	Not just the open file. Understands structure, git history, and how services talk to each other.
Plans before it acts	Every change shown as a diff, line by line, before anything is touched. No silent auto-apply.
Policy as code	CLAUDE.md is a plain-English policy file committed to the repo. Same rules for all 180 engineers.
Governance hook	Tests run automatically on every change. No passing tests, no commit.
Clean audit trail	Every tool call logged. Every commit attributed to the engineer, not the bot.

The numbers (conservative)

Team	Today	With Claude Code	Typical savings
Software Engineers (120) SREs (20)	~1 week to ramp on unfamiliar code ~45 min to find the issue at 2am	~1 day ~15 min	~50% of onboarding time ~60% of find-time
Data Scientists (40)	Blocked on boilerplate pipelines	Ship ML work, not plumbing	~70% of boilerplate hours

Unit math: \$1,500 in saved engineer-time per engineer per month vs. \$100 in software cost. 15× ratio. Assumes conservative adoption and uses FinTechCo’s fully-loaded hourly rate.

Security posture

- **SOC 2 Type II** certified. Report available on request.
- **Contractual no-training clause.** Not a policy, not a promise — a contract term your legal team will read before you sign.
- **PCI-DSS aware, FFIEC ready.** Enterprise agreements built with regulated industries in mind.
- **HIPAA eligible** under BAA. 200k-token context window. All traffic HTTPS to api.anthropic.com.
- **No source code uploaded in bulk.** Files are read locally; only the session prompt transits the network.

A pragmatic rollout

Phase	When	What	Cost
1 · Signal	Weeks 1-3	12 engineers across all three teams. Pick two real projects. Measure cycle time, incident MTTR, and engineer sentiment.	~\$3.6k
2 · Scale	Weeks 4-8	Expand to 40 engineers if phase 1 signal is positive. Publish CLAUDE.md policy. Roll out review plugin.	~\$12k

Phase	When	What	Cost
3 • Decide	Weeks 9-12	Readout with real data from your team. You decide whether to broaden or stop.	—

What we need from you to start

1. **This week** — security overview, DPA, and SOC 2 report to your legal team so their review can run in parallel.
2. **Next week** — a 30-minute setup call with the person you designate to own the pilot.
3. **Week 6** — a pilot readout meeting on your calendar.

No procurement commitment required today. No bulk license purchase required to run the pilot.

Honest limitations

Claude Code will not replace senior engineers, will not eliminate code review, and will not fix a broken engineering culture. It works best when there’s a real test suite to validate against. It gets things wrong sometimes — that’s exactly why the Safety Gate, governance hooks, and human approval exist.

Who’s your partner

Anthropic is an AI safety lab. Safety isn’t a feature we bolted on; it’s the training objective. That’s why Claude plans before it acts, why it shows its reasoning, and why your auditors will recognize the compliance posture when they read it.

Contact: Holden Ottolini · Solutions Engineering · Anthropic · holdenstirling@gmail.com